



Information Storage and Management I

Dr. Alejandro Arbelaez

- Labs commence in week 2 (next week - Thursdays) - **G.24 & G.21**
- Canvas is now working 😊



The MySQL Server

- A MySQL server can be given SQL commands, executes them, and results the results to the connected application:
- Starting the server
- Stopping the server

Command Line Interface

- Starting the MySQL client
- Commands for looking around in the database
 - SHOW DATABASES
 - USE database
 - SHOW TABLES
 - DESCRIBE table

Tables

- A typical database table definition has:
- A name
- A list of columns and their data types
- A list of constraints
 - Primary key to ensure uniqueness (if needed)
 - Foreign keys to facilitate relationships with other tables
 - Indices to facilitate fast look ups

Primary Key

- A **primary key** is a field in a table which uniquely identifies each row/record in a database table.
- **Primary keys** must contain unique values.
- A **primary key** column cannot have NULL values.
- A table can have only one **primary key**, which may consist of single or multiple fields.

employee_id	course_id	taken_date
100	3	1987-06-17
101	3	1989-09-21
102	3	1993-01-13
103	3	1990-01-03
104	3	1991-05-21
105	3	1997-06-25
106	3	1998-02-05

Primary Key

```
CREATE TABLE Customers (  
    ID INT NOT NULL,  
    NAME VARCHAR (20) NOT NULL,  
    AGE INT NOT NULL,  
    ADDRESS CHAR (25) ,  
    SALARY DECIMAL (18, 2)  
);
```

Primary Key

```
CREATE TABLE Customers (  
    ID INT NOT NULL,  
    NAME VARCHAR (20) NOT NULL,  
    AGE INT NOT NULL,  
    ADDRESS CHAR (25) ,  
    SALARY DECIMAL (18, 2)  
);
```

We are missing primary key(s)



```
ALTER TABLE Customers  
ADD PRIMARY KEY (ID);
```


Primary Key

```
CREATE TABLE Customers (  
    ID INT NOT NULL,  
    NAME VARCHAR (20) NOT NULL,  
    AGE INT NOT NULL,  
    ADDRESS CHAR (25) ,  
    SALARY DECIMAL (18, 2),  
    PRIMARY KEY (ID)  
);
```

Primary Key (with Multiple Columns)

```
CREATE TABLE Customers (  
    ID INT NOT NULL,  
    NAME VARCHAR (20) NOT NULL,  
    AGE INT NOT NULL,  
    ADDRESS CHAR (25) ,  
    SALARY DECIMAL (18, 2),  
    PRIMARY KEY (ID, NAME)  
);
```

Delete Primary Key

```
ALTER TABLE Customers  
DROP PRIMARY KEY ;
```

You can clear the primary key constraints from the table with the syntax given below.

Database structure for Quizzes

id	description	creation_date
1	Favorite Things Quiz	10/07/2014 10:22

quiz table

id	text	quiz_id
1	What is your favorite color?	1
2	What is your favorite book?	1
...		

question table

id	text	point_value	question_id
1	Red	1	1
2	Green	10	1
...			

answer table

Creating The Database Schema

```
CREATE TABLE quiz (  
    id INT NOT NULL AUTO_INCREMENT,  
    description VARCHAR(255),  
    create_time DATETIME NOT NULL,  
    PRIMARY KEY(id)  
)
```

```
CREATE TABLE question (  
    id INT NOT NULL AUTO_INCREMENT,  
    text VARCHAR(255),  
    quiz_id INT NOT NULL,  
    PRIMARY KEY (id),  
    FOREIGN KEY (quiz_id) REFERENCES  
    quiz(id) ON DELETE CASCADE  
)
```

```
CREATE TABLE answer (  
    id INT NOT NULL AUTO_INCREMENT,  
    text VARCHAR(255) NOT NULL,  
    point_value INT NOT NULL,  
    question_id INT NOT NULL,  
    PRIMARY KEY(id),  
    FOREIGN KEY (question_id) REFERENCES  
    question(id) ON DELETE CASCADE  
)
```

Filling In the Details

- Data types:
 - Numbers: INT, LONGINT, NUMERIC, FLOAT, DOUBLE
 - Strings: VARCHAR(<<NUM CHARS>>), TEXT, BLOB
 - Other: DATETIME
- NOT NULL vs NULL: whether to allow empty values or not
- PRIMARY KEY and FOREIGN KEY
- CASCADE: Keeping the data clean and robust

Null Values

- It is possible for tuples to have null values, denoted by null, for some of their attributes
- Null signifies an unknown value or that a value does not exist
- The result of any arithmetic expression involving null is null
 - Example: $5 + \text{null}$ returns null
- The predicate is null can be used to check for null values
 - **SELECT** name **FROM** instructor **WHERE** salary is null
- The predicate null is not null success if the value on which it is applied is not null

Null Values (Cont.)

- SQL treats as unknown the result of any comparison involving a null value (other than predicates is null and is not null).
 - Example: $5 < \text{null}$ or $\text{null} <> \text{null}$ or $\text{null} = \text{null}$
- The predicate in a where clause can involve Boolean operations (and, or, not); thus the definitions of the Boolean operations need to be extended to deal with the value unknown.
 - and : $(\text{true and unknown}) = \text{unknown}$,
 $(\text{false and unknown}) = \text{false}$,
 $(\text{unknown and unknown}) = \text{unknown}$
 - or: $(\text{unknown or true}) = \text{true}$,
 $(\text{unknown or false}) = \text{unknown}$
 $(\text{unknown or unknown}) = \text{unknown}$
- Result of where clause predicate is treated as false if it evaluates to unknown

Aggregate Functions

These functions operate on the multiset of values of a column of a relation, and return a value

- avg: average value

- min: minimum value

- max: maximum value

- sum: sum of values

- count: number of values

Aggregate Functions Examples

- Find the average salary of instructors in the Computer Science department
 - **SELECT AVG(salary) FROM instructor WHERE dept_name= 'Comp. Sci.'**
- Find the total number of instructors who teach a course in the Spring 2010 semester
 - **SELECT COUNT (distinct ID) FROM teaches WHERE semester = 'Spring' and year = 2018**
- Find the number of tuples in the course relation
 - **SELECT COUNT (*) FROM course;**

Aggregate Functions – Group By

- Find the average salary of instructors in each department

SELECT [?]

FROM instructor **GROUP BY** dept_name;

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

Aggregate Functions – Group By

- Find the average salary of instructors in each department

SELECT dept_name, AVG (salary) AS avg_salary **FROM** instructor
GROUP BY dept_name;

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

Aggregate Functions – Group By

- Find the average salary of instructors in each department

SELECT dept_name, AVG (salary) AS avg_salary **FROM** instructor
GROUP BY dept_name;

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

<i>dept_name</i>	<i>avg_salary</i>
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

Aggregation (Cont.)

Attributes in select clause outside of aggregate functions must appear in group by list

```
SELECT dept_name, ID, AVG (salary)  
FROM instructor  
GROUP BY dept_name;
```

correct or incorrect



Aggregation (Cont.)

Attributes in select clause outside of aggregate functions must appear in group by list

```
/* erroneous query */  
SELECT dept_name, ID, AVG (salary)  
FROM instructor  
GROUP BY dept_name;
```



Aggregate Functions – Having Clause

Find the names and average salaries of all departments whose average salary is greater than 42000

```
SELECT dept_name, AVG (salary) AS avg_salary  
FROM instructor  
GROUP BY dept_name  
HAVING AVG (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

Null Values and Aggregates

- Total all salaries

`SELECT SUM (salary) FROM instructor`

- Above statement ignores null amounts
 - Result is null if there is no non-null amount
- All aggregate operations except count(*) ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
 - count returns 0
 - all other aggregates return null

Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A subquery is a select-from-where expression that is nested within another query.
- The nesting can be done in the following SQL query

SELECT A1, A2, ..., An **FROM** r1, r2, ..., rm **WHERE** P

as follows:

From clause: r_i can be replaced by any valid subquery

Where clause: P can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

Where B is an attribute and $<\text{operation}>$ to be defined later.

Select clause:

A_i can be replaced by a subquery that generates a single value.

Set Membership

- Find courses offered in Fall 2017 and in Spring 2018

```
SELECT DISTINCT course_id
FROM section
WHERE semester = 'Fall' AND year= 2017 AND
       course_id IN (SELECT course_id
                       FROM section
                       WHERE semester = 'Spring' AND year= 2018);
```

- Find courses offered in Fall 2017 but not in Spring 2018

```
SELECT DISTINCT course_id
FROM section
WHERE semester = 'Fall' AND year= 2017 AND
       course_id NOT IN (SELECT course_id
                           FROM section
                           WHERE semester = 'Spring' AND year= 2018);
```

Set Membership (Cont.)

- Name all instructors whose name is neither “Mozart” nor Einstein”

```
SELECT distinct name FROM instructor  
WHERE name NOT IN ('Mozart', 'Einstein')
```

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
SELECT COUNT (distinct ID) FROM takes  
WHERE (course_id, sec_id, semester, year) IN  
      (SELECT course_id, sec_id, semester, year  
       FROM teaches  
       WHERE teaches.ID= 10101);
```

Set Comparison

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
SELECT DISTINCT T.name from instructor as T, instructor AS S  
WHERE T.salary > S.salary AND S.dept name = 'Biology';
```

Scalar Subquery

- Scalar subquery is one which is used where a single value is expected
- List all departments along with the number of instructors in each department

```
SELECT dept_name,  
      ( SELECT count(*)  
        FROM instructor  
        WHERE department.dept_name = instructor.dept_name)  
      AS num_instructors  
FROM department;
```

- Runtime error if subquery returns more than one result tuple

Modification of the Database

- Deletion of tuples from a given relation/table.
- Insertion of new tuples into a given relation/table
- Updating of values in some tuples in a given relation/table

Deletion

- Delete all instructors

DELETE FROM instructor

- Delete all instructors from the Finance department

DELETE FROM instructor
WHERE dept_name= 'Finance';

Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
DELETE FROM instructor  
WHERE salary < (SELECT AVG(salary)  
                  FROM instructor);
```

Problem: as we delete tuples from deposit, the average salary changes

Solution used in SQL:

1. First, compute **avg** (salary) and find all tuples to delete
2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

Insertion

- Add a new tuple to course

```
INSERT INTO course  
VALUES ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
INSERT INTO course (course_id, title, dept_name, credits)  
VALUES ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to student with tot_creds set to null

```
INSERT INTO student  
VALUES ('3003', 'Green', 'Finance', null);
```

Insertion (Cont.)

- Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000.

```
INSERT INTO instructor  
SELECT ID, name, dept_name, 18000  
FROM student  
WHERE dept_name = 'Music' and total_cred > 144;
```

- The select from where statement is evaluated fully before any of its results are inserted into the relation.
- Otherwise queries like

```
INSERT INTO table1 SELECT * FROM table1
```

would cause problem

Updates

- Give a 5% salary raise to all instructors

```
UPDATE instructor  
SET salary = salary * 1.05
```

- Give a 5% salary raise to those instructors who Eran less than 70000

```
UPDATE instructor  
SET salary = salary * 1.05  
WHERE salary < 70000;
```

- Give a 5% salary raise to instructors whose salary is less than average

```
UPDATE instructor  
SET salary = salary * 1.05  
WHERE salary < (select avg (salary)  
                  from instructor);
```

Updates (Cont.)

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
 - Write two update statements:
UPDATE instructor
SET salary = salary * 1.03
WHERE salary > 100000;
UPDATE instructor
SET salary = salary * 1.05
WHERE salary <= 100000;
 - The order is important
 - Can be done better using the case statement (next slide)

Case Statement for Conditional Updates

- Same query as before but with case statement

UPDATE instructor

SET salary = case

WHEN salary <= 100000 then salary * 1.05

ELSE salary * 1.03

END

Destroying and Altering Relations

DROP TABLE Students

- Destroys Students

ALTER TABLE Students **ADD COLUMN** firstYear

- Students will be altered by adding a new field
- Every tuple in the current instance is extended with a null value in the new field