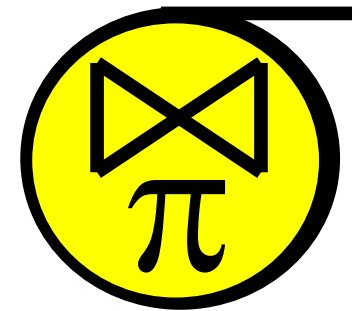




# Information Storage and Management I

Dr. Alejandro Arbelaez



Relational Algebra

- Labs commence this week (Thursdays 3-4 PM) – **1.10**
- **Continuous Assessment (20%)**
  - In-class test (10%) – October/17/2019
  - Project Assignment (10%) – November/28/2019



# Relational Query Languages

- **Query Languages**: Allow manipulation and **retrieval of data** from a database.
- **Relational model supports simple, powerful QLs**:
  - Strong **formal foundation** based on logic.
  - Allows for much optimization.

```
SELECT pname, price  
FROM Product, Company  
WHERE manufacturer=cname AND  
country = 'Japan' AND price < 150
```

No details of the  
implementation or how to get  
this efficiently get this data

# Relational Query Languages

- **Query Languages**: Allow manipulation and **retrieval of data** from a database.
- **Relational model supports simple, powerful QLs:**
  - Strong **formal foundation** based on logic.
  - Allows for much optimization.
- **Query Languages != programming languages!**
  - QLs not expected to be “Turing complete”.
  - QLs not intended to be used for complex calculations.
  - QLs support easy, efficient access to large data sets.

# Formal Relational Query Languages

Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:

- **Relational Algebra:** More *operational*, very useful for representing execution plans.
- **Relational Calculus:** Lets users describe what they want, rather than how to compute it. (*Non-operational, declarative.*)

*Understanding Algebra & Calculus is key to understanding SQL, query processing!*

# Preliminaries

- A query is applied to *relation instances*, and the result of a query is also a relation instance.
  - *Schemas of input* relations for a query are **fixed** (but query will run regardless of instance!)
  - The **schema for the result** of a given query is also **fixed**! Determined by definition of query language constructs.
- Positional vs. named-field notation:
  - Positional notation easier for formal definitions, named-field notation more readable.
  - Both used in Relational Algebra and SQL

# Example Instances

Database with the following relations

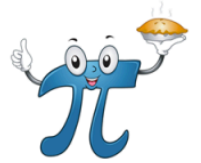
- R1(sid, bid, day)  
Key attributes: sid, bid, day  
Non-key attributes: None
- S1(sid, sname, rating, age)  
Key attributes: sid  
Non-key attributes: sname, rating, age
- S2(sid, sname, rating, age)  
Key attributes: sid  
Non-key attributes

Let's assume that names of fields in query results are ***inherited*** from names of fields in query input relations.

<i>R1</i>	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96

<i>S1</i>	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

<i>S2</i>	<u>sid</u>	sname	rating	age
	28	yuppy	9	35.0
	31	lubber	8	55.5
	44	guppy	5	35.0
	58	rusty	10	35.0



# Relational Algebra

- Basic operations:
  - **Selection** ( $\sigma$ ) Selects a subset of rows from relation.
  - **Projection** ( $\pi$ ) Deletes unwanted columns from relation.
  - **Cross-product** ( $\times$ ) Allows us to combine two relations.
  - **Set-difference** ( $-$ ) Tuples in reln. 1, but not in reln. 2.
  - **Union** ( $\cup$ ) Tuples in reln. 1 and in reln. 2.
  - **Renaming** ( $\rho$ ) (for named perspective)
- Additional operations:
  - Intersection, **join**, division, renaming: Not essential, but (very!) useful.
- Since each operation returns a relation, **operations can be composed!**



## Projection -- $\pi$

- Deletes attributes that are not in *projection list*.
- **Schema** of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- Projection operator has to eliminate **duplicates**! (Why??)
  - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\pi_{sname, rating}(S2)$

# Projection

- Deletes attributes that are not in *projection list*.
- **Schema** of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- Projection operator has to eliminate **duplicates**! (Why??)
  - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\pi_{sname, rating}(s2)$



sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

# Projection

- Deletes attributes that are not in *projection list*.
- **Schema** of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- Projection operator has to eliminate **duplicates**! (Why??)
  - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$$\pi_{age}(S2)$$

# Projection

- Deletes attributes that are not in *projection list*.
- **Schema** of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- Projection operator has to eliminate **duplicates**! (Why??)
  - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\pi_{age}(S2)$



age
35.0
55.5

## Selection -- $\sigma$

- Selects rows that satisfy *selection condition*
- No duplicates in result! (Why?)
- *Schema* of result identical to schema of (only) input relation.
- *Result* relation can be the *input* for another relational algebra operation! (*Operator composition.*)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

# Selection

- Selects rows that satisfy *selection condition*
- No duplicates in result! (Why?)
- *Schema* of result identical to schema of (only) input relation.
- *Result* relation can be the *input* for another relational algebra operation! (*Operator composition.*)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\sigma_{rating > 8}(S2)$



Same attributes



sid	sname	rating	age
-----	-------	--------	-----

# Selection

- Selects rows that satisfy *selection condition*
- No duplicates in result! (Why?)
- *Schema* of result identical to schema of (only) input relation.
- *Result* relation can be the *input* for another relational algebra operation! (*Operator composition.*)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\sigma_{rating > 8}(S2)$



sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

# Selection

- Selects rows that satisfy *selection condition*
- No duplicates in result! (Why?)
- *Schema* of result identical to schema of (only) input relation.
- *Result* relation can be the *input* for another relational algebra operation! (*Operator composition.*)

s2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$$



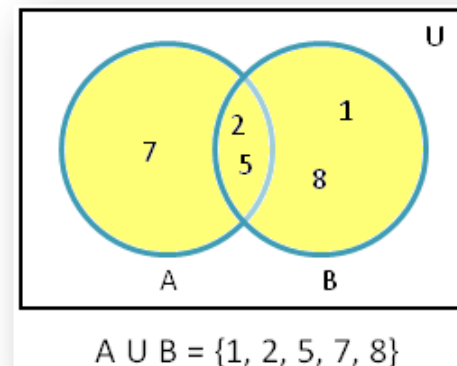
sname	rating
yuppy	9
rusty	10



## Union -- $\cup$

- Produces a resulting relation that contains a tuple for every tuple in either or both of two input relations (duplicates only occur once)
- The Relations being combined must be **union-compatible** (type-compatible)

Set  
theory



# Union

- All of these operations take two input relations, which must be *compatible*:
  - Same number of fields.
  - “Corresponding” fields have the same type.
- What is the *schema* of result?

$S1 \cup S2 \rightarrow$  sid, sname, rating , age

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*S2*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

# Union

- All of these operations take two input relations, which must be **compatible**:
  - Same number of fields.
  - “Corresponding” fields have the same type.
- What is the **schema** of result?

$S1 \cup S2 \rightarrow$  sid, sname, rating, age

**S1**

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**S2**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 \cup S2$

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

# Union

**SELECT \* FROM S1**

**UNION**

**SELECT \* FROM S2**

But be careful with duplicates

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*S2*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 \cup S2$

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

# Implementation

- $i=0$
- $j=0$
- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

# Implementation

- $i=0$
- $j=0$
- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union = { }

Two-pass algorithm based on  
Sorted  
**Sort First**

# Implementation

- $i=0$
- $j=0$
- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union = { }

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union = { }



# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union = {1}

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3, 4\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3, 4\}$



# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3, 4, 5\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3, 4, 5, 6\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3, 4, 5, 6\}$

# Implementation

- if  $T1[i] < T2[j]$  then print  $T1[i]$  and  $i+=1$
- elif  $T1[i] > T2[j]$  then print  $T2[j]$  and  $j+=1$
- elif both are the same then print any of them and increment both  $i$  and  $j$
- print the remaining elements of the larger array

$T1 = \{1, 3, 4, 5, 7\}$

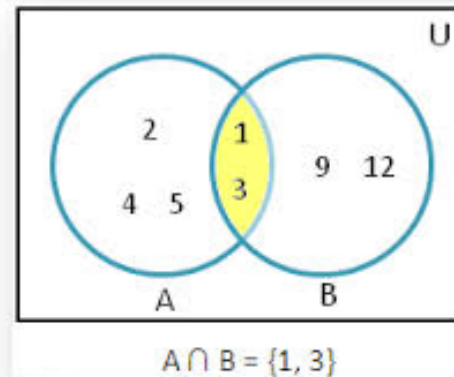
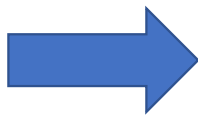
$T2 = \{2, 3, 5, 6\}$

Union =  $\{1, 2, 3, 4, 5, 6, 7\}$

# Intersection

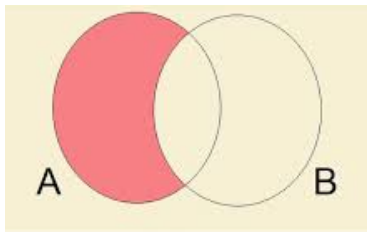
- Produces a resulting relation that contains a tuple for every tuple in BOTH of the two input relations
- The relations being combined must be **compatible** (type-compatible)

Set  
theory



## Set-Difference (Minus)

- Produces a resulting relation that contains a tuple for every tuple in the first of two input relations and not in the second.
- The Relations being combined must be union-compatible (type-compatible)



$$A = \{1, 2, 3, 4\} \quad B = \{2, 4, 6\}$$

$$A - B = \{1, 2, 3, 4\} - \{2, 4, 6\} = \{1, 3\}$$

$$B - A = \{2, 4, 6\} - \{1, 2, 3, 4\} = \{6\}$$

# Set-Difference (Minus)

- All of these operations take two input relations, which must be *compatible*:
  - Same number of fields.
  - “Corresponding” fields have the same type.
- What is the *schema* of result?

$S1 - S2 \rightarrow$  sid, sname, rating, age

*S1*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

*S2*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

# Set-Difference (Minus)

- All of these operations take two input relations, which must be *compatible*:
  - Same number of fields.
  - “Corresponding” fields have the same type.
- What is the *schema* of result?

$S1 - S2 \rightarrow$  sid, sname, rating, age

$S1$

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

$S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 - S2$

sid	sname	rating	age
-----	-------	--------	-----



# Set-Difference (Minus)

- All of these operations take two input relations, which must be *compatible*:
  - Same number of fields.
  - “Corresponding” fields have the same type.
- What is the *schema* of result?

$S1 - S2 \rightarrow$  sid, sname, rating, age

$S1$

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

$S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 - S2$

sid	sname	rating	age
22	dustin	7	45.0

# Set-Difference (Minus)

- All of these operations take two input relations, which must be *compatible*:
  - Same number of fields.
  - “Corresponding” fields have the same type.
- What is the *schema* of result?

$S1 - S2 \rightarrow$  sid, sname, rating, age

Unfortunately, MySQL doesn't support MINUS operator

$S1$

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

$S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$S1 - S2$

sid	sname	rating	age
22	dustin	7	45.0